

Analysis of Algorithms, I

CSOR W4231.002

Eleni Drinea
Computer Science Department

Columbia University

Tuesday, January 26, 2016

Outline

- 1 Recap
- 2 Binary search
- 3 Integer multiplication
- 4 Fast matrix multiplication
- 5 Quicksort

Today

1 Recap

2 Binary search

3 Integer multiplication

4 Fast matrix multiplication

5 Quicksort

Review of the last lecture

- ▶ Asymptotic notation ($O, \Omega, \Theta, o, \omega$)
- ▶ The divide & conquer principle
 - ▶ **Divide** the problem into a number of subproblems that are smaller instances of the same problem.
 - ▶ **Conquer** the subproblems by solving them recursively.
 - ▶ **Combine** the solutions to the subproblems into the solution for the original problem.
- ▶ Application: **Mergesort**
- ▶ Solving recurrences

Mergesort

```
Mergesort (A, left, right)  
  if right == left then  
    return  
  end if  
  middle = left +  $\lfloor (\textit{right} - \textit{left}) / 2 \rfloor$   
  Mergesort (A, left, middle)  
  Mergesort (A, middle + 1, right)  
  Merge (A, left, middle, right)
```

- ▶ Initial call: Mergesort(*A*, 1, *n*)
- ▶ Subroutine Merge merges two **sorted** lists of sizes $\lceil n/2 \rceil$, $\lfloor n/2 \rfloor$ into one sorted list of size *n* in time $\Theta(n)$.

Running time of Mergesort

The running time of **Mergesort** satisfies:

$$T(n) = 2T(n/2) + cn, \text{ for } n \geq 2, \text{ constant } c > 0$$

$$T(1) = c$$

This structure is typical of **recurrence relations**:

- ▶ an inequality or equation bounds $T(n)$ in terms of an expression involving $T(m)$ for $m < n$
- ▶ a base case generally says that $T(n)$ is constant for small constant n

Remarks

- ▶ We ignore floor and ceiling notations
- ▶ A recurrence does **not** provide an asymptotic bound for $T(n)$: to this end, we must **solve** the recurrence

Solving recurrences, method 1: recursion trees

The technique consists of three steps

1. Analyze the first few levels of the tree of recursive calls
2. Identify a pattern
3. Sum over all levels of recursion

Example: analysis of running time of Mergesort

$$T(n) = 2T(n/2) + cn, n \geq 2$$

$$T(1) = c$$

A general recurrence and its solution

Often we can express the running time of a recursive algorithm by the following recurrence

$$T(n) = aT(n/b) + cn^k, \text{ for } a, c > 0, b > 1, k \geq 0$$

What is the recursion tree for this recurrence?

- ▶ a is the branching factor
- ▶ b is the factor by which the size of each subproblem shrinks
- ⇒ at level i , there are a^i subproblems, each of size n/b^i
- ⇒ each subproblem at level i requires $c(n/b^i)^k$ work
- ▶ the height of the tree is $\log_b n$ levels
- ⇒ Total work: $\sum_{i=0}^{\log_b n} a^i c(n/b^i)^k = cn^k \sum_{i=0}^{\log_b n} \left(\frac{a}{b^k}\right)^i$

Solving recurrences, method 2: Master theorem

Theorem 1 (Master theorem).

If $T(n) = aT(\lceil n/b \rceil) + O(n^k)$ for some constants $a > 0$, $b > 1$, $k \geq 0$, then

$$T(n) = \begin{cases} O(n^{\log_b a}) & , \text{ if } a > b^k \\ O(n^k \log n) & , \text{ if } a = b^k \\ O(n^k) & , \text{ if } a < b^k \end{cases}$$

Example: running time of Mergesort

- $T(n) = 2T(n/2) + cn$:
 $a = 2, b = 2, k = 1, b^k = 2 = a \Rightarrow T(n) = O(n \log n)$

Today

1 Recap

2 Binary search

3 Integer multiplication

4 Fast matrix multiplication

5 Quicksort

Searching a sorted array

- ▶ **Input:** sorted list A of n integers, integer x
- ▶ **Output:**
 1. index j s.t. $1 \leq j \leq n$ and $A[j] = x$; or
 2. **no** if x is not in A

Searching a sorted array

- ▶ **Input:** sorted list A of n integers, integer x
- ▶ **Output:**
 1. index j s.t. $1 \leq j \leq n$ and $A[j] = x$; or
 2. **no** if x is not in A

Example: $A = \{0, 2, 3, 5, 6, 7, 9, 11, 13\}$, $n = 9$, $x = 7$

Searching a sorted array

- ▶ **Input:** sorted list A of n integers, integer x
- ▶ **Output:**
 1. index j s.t. $1 \leq j \leq n$ and $A[j] = x$; or
 2. **no** if x is not in A

Example: $A = \{0, 2, 3, 5, 6, 7, 9, 11, 13\}$, $n = 9$, $x = 7$

Idea: use the fact that the array is **sorted** and probe specific entries in the array.

Binary search

First, probe the middle entry. Let $mid = \lceil n/2 \rceil$.


- ▶ If $x == A[mid]$, return mid .
- ▶ If $x < A[mid]$ then look for x in $A[1, mid - 1]$;
- ▶ Else if $x > A[mid]$ look for x in $A[mid + 1, n]$.

Initially, the entire array is “active”, that is, x might be anywhere in the array.



Suppose $x > A[mid]$.

Then the active area of the array, where x might be, is to the right of mid .



Binary search pseudocode

```
binarysearch(A, left, right)  
   $mid = left + \lceil (right - left) / 2 \rceil$   
  if  $A[mid] == x$  then  
    return mid  
  else if  $right == left$  then  
    return no  
  else if  $A[mid] < x$  then  
     $left = mid + 1$   
  else  $right = mid - 1$   
  end if  
  binarysearch(A, left, right)
```

Binary search running time

Observation: At each step there is a region of A where x could be and we **shrink** the size of this region by a factor of 2 with every probe:

- ▶ If n is odd, then we are throwing away $\lceil n/2 \rceil$ elements.
- ▶ If n is even, then we are throwing away at least $n/2$ elements.

Binary search running time

Observation: At each step there is a region of A where x could be and we **shrink** the size of this region by a factor of 2 with every probe:

- ▶ If n is odd, then we are throwing away $\lceil n/2 \rceil$ elements.
- ▶ If n is even, then we are throwing away at least $n/2$ elements.

Hence the recurrence for the running time is

$$T(n) \leq T(n/2) + O(1)$$

Sublinear running time

Here are two ways to argue about the running time:

1. Master theorem: $b = 2, a = 1, k = 0 \Rightarrow T(n) = O(\log n)$.
2. We can reason as follows:
 - ▶ Starting with an array of size n , after k probes, we are left with an array of size at most $\frac{n}{2^k}$ (since every time we probe an entry the active portion of the array halves).
 - ▶ Hence after $k = \log n$ probes, we are left with an array of **constant** size (i.e., $O(1)$). Now we can search **linearly** for x in the constant size array.

Concluding remarks on binary search

1. The right data structure can improve the running time of the algorithm significantly.
 - ▶ *What if we used a **linked list** to store the input?*
 - ▶ Arrays allow for **random access** of their elements: given an index, we can read any entry in an array in time $O(1)$ (constant time)
2. In general, we obtain running time $O(\log n)$ when the algorithm does a **constant amount of work** to throw away a **constant fraction** of the input.

Today

- 1 Recap
- 2 Binary search
- 3 Integer multiplication**
- 4 Fast matrix multiplication
- 5 Quicksort

Integer multiplication

- ▶ *How do we multiply two integers x and y ?*
- ▶ Elementary school method: compute a partial product by multiplying every digit of y separately with x and then add up all the partial products.
- ▶ Remark: this method works the same in base 10 or base 2.

Examples: $(12)_{10} \cdot (11)_{10}$ and $(1100)_2 \cdot (1011)_2$

$$\begin{array}{r} 12 \\ \times 11 \\ \hline 12 \\ + 12 \\ \hline 132 \end{array}$$

$$\begin{array}{r} 1100 \\ \times 1011 \\ \hline 1100 \\ 1100 \\ 0000 \\ + 1100 \\ \hline 10000100 \end{array}$$

Elementary algorithm running time

A more reasonable model of computation: a **single** operation on a pair of digits (bits) is a primitive computational step.

Assume we are multiplying n -digit (bit) numbers.

- ▶ $O(n)$ time to compute a partial product.
 - ▶ $O(n)$ time to combine it in a running sum of all partial products so far.
- ⇒ There are n partial products, each consisting of n bits, hence total number of operations is $O(n^2)$.

Can we do better?

A first divide & conquer approach

Consider n -digit decimal numbers x, y .

$$x = x_{n-1}x_{n-2} \dots x_0$$

$$y = y_{n-1}y_{n-2} \dots y_0$$

Idea: rewrite each number as the sum of the $n/2$ high-order digits and the $n/2$ low-order digits.

$$x = \underbrace{x_{n-1} \dots x_{n/2}}_{x_H} \underbrace{x_{n/2-1} \dots x_0}_{x_L} = x_H \cdot 10^{n/2} + x_L$$

$$y = \underbrace{y_{n-1} \dots y_{n/2}}_{y_H} \underbrace{y_{n/2-1} \dots y_0}_{y_L} = y_H \cdot 10^{n/2} + y_L$$

where each of x_H, x_L, y_H, y_L is an $n/2$ -digit number.

Examples

- $n = 2, x = 12, y = 11$

$$\underbrace{12}_x = \underbrace{1}_{x_H} \cdot \underbrace{10^1}_{10^{n/2}} + \underbrace{2}_{x_L}$$

$$\underbrace{11}_y = \underbrace{1}_{y_H} \cdot \underbrace{10^1}_{10^{n/2}} + \underbrace{1}_{y_L}$$

- $n = 4, x = 1000, y = 1110$

$$\underbrace{1000}_x = \underbrace{10}_{x_H} \cdot \underbrace{10^2}_{10^{n/2}} + \underbrace{0}_{x_L}$$

$$\underbrace{1110}_y = \underbrace{11}_{y_H} \cdot \underbrace{10^2}_{10^{n/2}} + \underbrace{10}_{y_L}$$

A first divide & conquer approach

$$\begin{aligned}x \cdot y &= (x_H 10^{n/2} + x_L) \cdot (y_H 10^{n/2} + y_L) \\&= x_H y_H \cdot 10^n + (x_H y_L + x_L y_H) 10^{n/2} + x_L y_L\end{aligned}$$

In words, we reduced the problem of solving 1 instance of size n (i.e., one multiplication between two n -digit numbers) to the problem of solving 4 instances, each of size $n/2$ (i.e., computing the products $x_H y_H$, $x_H y_L$, $x_L y_H$ and $x_L y_L$).

A first divide & conquer approach

$$\begin{aligned}x \cdot y &= (x_H 10^{n/2} + x_L) \cdot (y_H 10^{n/2} + y_L) \\&= x_H y_H \cdot 10^n + (x_H y_L + x_L y_H) 10^{n/2} + x_L y_L\end{aligned}$$

In words, we reduced the problem of solving 1 instance of size n (i.e., one multiplication between two n -digit numbers) to the problem of solving 4 instances, each of size $n/2$ (i.e., computing the products $x_H y_H$, $x_H y_L$, $x_L y_H$ and $x_L y_L$).

This is a **divide and conquer** solution!

- ▶ Recursively solve the 4 subproblems.
- ▶ Multiplication by 10^n is easy (**shifting**): $O(n)$ time.
- ▶ Combine the solutions from the 4 subproblems to an overall solution using 3 additions on $O(n)$ -digit numbers: $O(n)$ time.

Karatsuba's observation

Running time: $T(n) \leq 4T(n/2) + cn$

- ▶ by the Master Theorem: $T(n) = O(n^2)$
- ▶ **no** improvement

Karatsuba's observation

Running time: $T(n) \leq 4T(n/2) + cn$

- ▶ by the Master Theorem: $T(n) = O(n^2)$
- ▶ **no** improvement

However, **if we only needed three $n/2$ -digit multiplications**, then by the Master theorem

$$T(n) \leq 3T(n/2) + cn = O(n^{1.59}) = o(n^2).$$

Karatsuba's observation

Running time: $T(n) \leq 4T(n/2) + cn$

- ▶ by the Master Theorem: $T(n) = O(n^2)$
- ▶ **no** improvement

However, **if we only needed three $n/2$ -digit multiplications**, then by the Master theorem

$$T(n) \leq 3T(n/2) + cn = O(n^{1.59}) = o(n^2).$$

Recall that

$$x \cdot y = x_H y_H \cdot 10^n + (x_H y_L + x_L y_H) 10^{n/2} + x_L y_L$$

Key observation: we do not need **each** of $x_H y_L, x_L y_H$.
We only need **their sum**, $x_H y_L + x_L y_H$.

Gauss's observation on multiplying complex numbers

A similar problem: multiply two complex numbers $a + bi, c + di$

$$(a + bi)(c + di) = ac + (ad + bc)i + bdi^2$$

Gauss's observation on multiplying complex numbers

A similar problem: multiply two complex numbers $a + bi, c + di$

$$(a + bi)(c + di) = ac + (ad + bc)i + bdi^2$$

Gauss's observation: can be done with just 3 multiplications

$$(a + bi)(c + di) = ac + ((a + b)(c + d) - ac - bd)i + bdi^2,$$

at the cost of few extra additions and subtractions.

* Unlike multiplications, additions and subtractions of n -digit numbers are cheap: $O(n)$ time!

Karatsuba's algorithm

$$\begin{aligned}x \cdot y &= (x_H 10^{n/2} + x_L) \cdot (y_H 10^{n/2} + y_L) \\&= x_H y_H \cdot 10^n + (x_H y_L + x_L y_H) 10^{n/2} + x_L y_L\end{aligned}$$

Similarly to Gauss's method for multiplying two complex numbers, compute only the three products

$$x_H y_H, x_L y_L, (x_H + x_L)(y_H + y_L)$$

and obtain the sum $x_H y_L + x_L y_H$ from

$$(x_H + x_L)(y_H + y_L) - x_H y_H - x_L y_L = x_H y_L + x_L y_H.$$

Combining requires $O(n)$ time hence

$$T(n) \leq 3T(n/2) + cn = O(n^{\log_2 3}) = O(n^{1.59})$$

Pseudocode

Let k be a small constant.

Integer-Multiply(x, y)

if $n == k$ **then**

 return xy

end if

 write $x = x_H 10^{n/2} + x_L, y = y_H 10^{n/2} + y_L$

 compute $x_H + x_L, y_H + y_L$

$product = \text{Integer-Multiply}(x_H + x_L, y_H + y_L)$

$x_H y_H = \text{Integer-Multiply}(x_H, y_H)$

$x_L y_L = \text{Integer-Multiply}(x_L, y_L)$

 return $x_H y_H 10^n + (product - x_H y_H - x_L y_L) 10^{n/2} + x_L y_L$

Concluding remarks

- ▶ To reduce the number of multiplications we do few more additions/subtractions: these are fast compared to multiplications.
- ▶ There is no reason to continue with recursion once n is small enough: the conventional algorithm is probably more efficient since it uses fewer additions.
- ▶ When we recursively compute $(x_H + x_L)(y_H + y_L)$, each of $x_H + x_L$, $y_H + y_L$ might be $(n/2 + 1)$ -digit integers. This does not affect the asymptotics.

Today

- 1 Recap
- 2 Binary search
- 3 Integer multiplication
- 4 Fast matrix multiplication**
- 5 Quicksort

Fast matrix multiplication

Matrix multiplication: a fundamental primitive in numerical linear algebra, scientific computing, machine learning and large-scale data analysis.

- ▶ Input: $m \times n$ matrix A , $n \times p$ matrix B
- ▶ Output: $m \times p$ matrix $C = AB$

Example: $A = \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix}, B = \begin{pmatrix} 1 & 1 \\ 1 & 1 \end{pmatrix}, C = \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix}$

Lower bounds on matrix multiplication algorithms for $m, p = \Theta(n)$?

Conventional matrix multiplication

```
for  $1 \leq i \leq m$  do  
  for  $1 \leq j \leq p$  do  
     $c_{i,j} = 0$   
    for  $1 \leq k \leq n$  do  
       $c_{i,j} + = a_{i,k} \cdot b_{k,j}$   
    end for  
  end for  
end for
```

- ▶ *Running time?*
- ▶ *Can we do better?*

A first divide & conquer approach: 8 subproblems

Assume square A, B where $n = 2^k$ for some $k > 0$.

Idea: express A, B as 2×2 block matrices and use the conventional algorithm to multiply the two block matrices.

$$\begin{pmatrix} \overbrace{A_{11} \quad A_{12}}^{n/2 \times n/2} \\ A_{21} \quad A_{22} \end{pmatrix} \begin{pmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{pmatrix} = \begin{pmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{pmatrix}$$

where

$$C_{11} = A_{11}B_{11} + A_{12}B_{21}$$

$$C_{12} = A_{11}B_{12} + A_{12}B_{22}$$

$$C_{21} = A_{21}B_{11} + A_{22}B_{21}$$

$$C_{22} = A_{21}B_{12} + A_{22}B_{22}$$

Running time?

Strassen's breakthrough: 7 subproblems suffice (part 1)

Compute the following ten $n/2 \times n/2$ matrices.

1. $S_1 = B_{11} - B_{22}$
2. $S_2 = A_{11} + A_{12}$
3. $S_3 = A_{21} + A_{22}$
4. $S_4 = B_{21} - B_{11}$
5. $S_5 = A_{11} + A_{22}$
6. $S_6 = B_{11} + B_{22}$
7. $S_7 = A_{12} - A_{22}$
8. $S_8 = B_{21} + B_{22}$
9. $S_9 = A_{11} - A_{21}$
10. $S_{10} = B_{11} + B_{12}$

Running time?

Strassen's breakthrough: 7 subproblems suffice (part 2)

Compute the following seven products of $n/2 \times n/2$ matrices.

1. $P_1 = A_{11}S_1$

2. $P_2 = S_2B_{22}$

3. $P_3 = S_3B_{11}$

4. $P_4 = A_{22}S_4$

5. $P_5 = S_5S_6$

6. $P_6 = S_7S_8$

7. $P_7 = S_9S_{10}$

Compute C as follows:

1. $C_{11} = P_4 + P_5 + P_6 - P_2$

2. $C_{12} = P_1 + P_2$

3. $C_{21} = P_3 + P_4$

4. $C_{22} = P_1 + P_5 - P_3 - P_7$

Running time?

Strassen's running time and concluding remarks

- ▶ Recurrence: $T(n) = 7T(n/2) + cn^2$
- ▶ By the Master theorem:

$$T(n) = O(n^{\log_2 7}) = O(n^{2.81})$$

- ▶ Recently, there is renewed interest in Strassen's algorithm for **high-performance computing**: thanks to its lower communication cost (number of bits exchanged between machines in the network or data center), it is better suited than the traditional algorithm for multi-core processors.

Today


- 1 Recap
- 2 Binary search
- 3 Integer multiplication
- 4 Fast matrix multiplication
- 5 Quicksort**

Quicksort facts

- ▶ Quicksort is a **divide and conquer** algorithm
- ▶ It is the standard algorithm used for sorting
- ▶ It is an **in-place** algorithm
- ▶ Its worst-case running time is $\Theta(n^2)$ but its average-case running time is $\Theta(n \log n)$
- ▶ We will use it to introduce **randomized** algorithms

Quicksort: main idea

- ▶ Pick an input item, call it *pivot*, and place it in its **final location** in the sorted array by **re-organizing** the array:
 - ▶ all items \leq *pivot* are placed **before** *pivot*
 - ▶ all items $>$ *pivot* are placed **after** *pivot*



- ▶ Recursively sort the subarray to the left of *pivot* (items \leq *pivot*).
- ▶ Recursively sort the subarray to the right of *pivot* (items $>$ *pivot*).

Remark: I haven't explained how to pick *pivot*.

Quicksort pseudocode

```
Quicksort( $A, left, right$ )  
  if  $|A| = 0$  then return           //  $A$  is empty  
  end if  
   $split = \text{Partition}(A, left, right)$   
  Quicksort( $A, left, split - 1$ )  
  Quicksort( $A, split + 1, right$ )
```

Initial call: Quicksort($A, 1, n$)


Subroutine Partition($A, left, right$)

Notation: $A[i, j]$ denotes the portion of A starting at position i and ending at position j .

Partition

1. **picks** a *pivot* item
2. **re-organizes** $A[left, right]$ so that
 - ▶ all items **before** *pivot* are \leq *pivot*
 - ▶ all items **after** *pivot* are $>$ *pivot*
3. returns *split*, the **index** of *pivot* in the re-organized array

After Partition, $A[left, right]$ looks as follows:



Implementing Partition


1. Pick a *pivot* item: for simplicity, always pick the **last item** of the array as *pivot*, i.e., $\text{pivot} = A[\text{right}]$.
 - ▶ Thus $A[\text{right}]$ will be placed in its final location when `Partition` returns and **will never be used (or moved) again until the algorithm terminates.**
2. Re-organize the input array A **in place**.

(What if we didn't care to implement `Partition` in place?)

Implementing Partition in place

Partition examines the items in $A[\text{left}, \text{right}]$ one by one and maintains **three regions** in A . Specifically, **after** examining the j -th item for $j \in [\text{left}, \text{right} - 1]$, the regions are:

1. **First region**: starts at $A[\text{left}]$ and ends at $A[\text{split}]$; it contains all items examined so far that are $\leq \text{pivot}$.
2. **Second region**: starts at $A[\text{split} + 1]$ and ends at $A[j]$; it contains all items $> \text{pivot}$ examined so far;
3. **Third region**: starts at $A[j + 1]$ and ends at $A[\text{right} - 1]$; it contains all items **yet to be examined**.



Implementing Partition in place (cont'd)


At the beginning of iteration j , $A[j]$ is compared with $pivot$.

If $A[j] \leq pivot$


1. swap $A[j]$ with $A[split + 1]$, the first element of the **second region**: since $A[split + 1] > pivot$, it's “safe” to move it to the end of the second region (recall that the second region ends at position j after examining the j -th item);
2. increment $split$ to include $A[j]$ in the **first region**, where all items are $\leq pivot$ (recall that the first region ends at $split$).

Iteration j , when $A[j] \leq pivot$

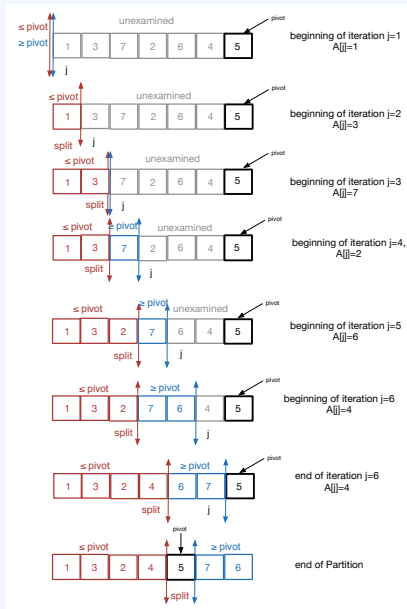
Start of iteration j (assume $A[j] \leq pivot$)



End of iteration j : $A[j]$ swapped with $A[split+1]$, $split$ updated to $split+1$



Example: $A = \{1, 3, 7, 2, 6, 4, 5\}$, Partition($A, 1, 7$)



Pseudocode for Partition

Partition($A, left, right$)

$pivot = A[right]$

$split = left - 1$

for $j = left$ to $right - 1$ **do**

if $A[j] \leq pivot$ **then**

$swap(A[j], A[split + 1])$

$split = split + 1$

end if

end for

$swap(pivot, A[split + 1])$ //place $pivot$ after $A[split]$ (*why?*)

$return\ split + 1$ //the final position of $pivot$

Analysis of Partition: correctness

Notation: $A[a, \dots, b]$ denotes the portion of A starting at position a and ending at position b .

Claim 1.

For $left \leq j \leq right - 1$, at the end of loop j , all items from $A[left, j]$ that appear in

- 1. $A[left, split]$ are less or equal to pivot*
- 2. $A[split + 1, j]$ are greater than pivot*

Remark: If the claim is true, correctness of **Partition** follows (*why?*).

Proof of Claim 1

By induction on j .

1. **Base case:** For $j = left$ (that is, during the first execution of the for loop), there are two possibilities:
 - ▶ if $A[left] \leq pivot$, then $A[left]$ is swapped with itself and $split$ is incremented to equal $left$;
 - ▶ otherwise, nothing happens.

In both cases, the claim holds for $j = left$.

2. **Hypothesis:** Assume that the claim is true for some $left \leq j < right - 1$.
 - ▶ That is, at the end of loop j , all items in $A[left, split]$ are $\leq pivot$ and all items in $A[split + 1, j]$ are $> pivot$.

Proof of Claim 1 (cont'd)

3. Step: We will show the claim for $j + 1$.

Thus we will show that after loop $j + 1$, all items in $A[\textit{left}, \textit{split}]$ are $\leq \textit{pivot}$ and all items in $A[\textit{split} + 1, j + 1]$ are $> \textit{pivot}$.

- ▶ At the beginning of loop $j + 1$, by the hypothesis, items in $A[\textit{left}, \textit{split}]$ are $\leq \textit{pivot}$ and items in $A[\textit{split} + 1, j]$ are $> \textit{pivot}$.
- ▶ Inside loop $j + 1$, there are two possibilities:
 1. $A[j + 1] \leq \textit{pivot}$: then $A[j + 1]$ is swapped with $A[\textit{split} + 1]$. At this point, items in $A[\textit{left}, \textit{split} + 1]$ are $\leq \textit{pivot}$ and items in $A[\textit{split} + 2, j + 1]$ are $> \textit{pivot}$. Incrementing \textit{split} (the next step in the pseudocode) yields that the claim holds for $j + 1$.
 2. $A[j + 1] > \textit{pivot}$: nothing is done. The truth of the claim follows from the hypothesis.

This completes the proof of the inductive step.

Analysis of Partition: running time and space

- ▶ **Running time:** on input size n , **Partition** goes through each of the $n - 1$ leftmost elements once and performs constant amount of work per element
 - ⇒ **Partition** requires $\Theta(n)$ time.
- ▶ **Space:** in-place algorithm