

# Analysis of Algorithms, I

## CSOR S4231

Eleni Drinea  
*Computer Science Department*

Columbia University

Thursday, February 4, 2016

# Outline

- 1 Recap
- 2 Data compression
- 3 Symbol codes and optimal lossless compression
- 4 Prefix codes
- 5 Prefix codes and trees
- 6 The Huffman algorithm

# Today

- 1 Recap
- 2 Data compression
- 3 Symbol codes and optimal lossless compression
- 4 Prefix codes
- 5 Prefix codes and trees
- 6 The Huffman algorithm

# Review of the last lecture

- ▶ Expected running time of `Randomized-Quicksort`
- ▶ Balls-in-bins problems
  - ▶ Birthday paradox
  - ▶ Expected load of a bin
  - ▶ Expected number of empty bins
  - ▶ Expected number of balls thrown before every bin has at least one ball

# Today

- 1 Recap
- 2 Data compression
- 3 Symbol codes and optimal lossless compression
- 4 Prefix codes
- 5 Prefix codes and trees
- 6 The Huffman algorithm

**Data compression:** find compact representations of data

Data compression **standards**

- ▶ jpeg for image transmission
- ▶ mp3 for audio content, mpeg2 for video transmission
- ▶ utilities: gzip, bzip2

All of the above use the **Huffman algorithm** as a basic building block.

# Data representation

- ▶ *Chromosome maps*: sequences of hundreds of millions of bases (symbols from  $\{A, C, G, T\}$ ).
- ▶ **Goal**: store a chromosome map with 200 million bases.

*How do we represent a chromosome map?*

# Data representation

- ▶ *Chromosome maps*: sequences of hundreds of millions of bases (symbols from  $\{A, C, G, T\}$ ).
- ▶ **Goal**: store a chromosome map with 200 million bases.

*How do we represent a chromosome map?*

- ▶ **Encode** every symbol that appears in the sequence separately by a **fixed length binary string**.
- ▶ **Codeword**  $c(x)$  for symbol  $x$ : a binary string encoding  $x$



## Example code

- ▶ Alphabet  $\mathcal{A} = \{A, C, G, T\}$  with 4 symbols
- ▶ Encode each symbol with 2 bits

alphabet symbol $x$	codeword $c(x)$
$A$	00
$C$	01
$G$	10
$T$	11

## Example code

- ▶ Alphabet  $\mathcal{A} = \{A, C, G, T\}$  with 4 symbols
- ▶ Encode each symbol with 2 bits

alphabet symbol $x$	codeword $c(x)$
$A$	00
$C$	01
$G$	10
$T$	11

**Output of encoding:** the concatenation of the codewords for every symbol in the input sequence

## Example code

- ▶ Alphabet  $\mathcal{A} = \{A, C, G, T\}$  with 4 symbols
- ▶ Encode each symbol with 2 bits

alphabet symbol $x$	codeword $c(x)$
$A$	00
$C$	01
$G$	10
$T$	11

**Output of encoding:** the concatenation of the codewords for every symbol in the input sequence

### Example

- ▶ Input sequence:  $ACGTAA$

## Example code

- ▶ Alphabet  $\mathcal{A} = \{A, C, G, T\}$  with 4 symbols
- ▶ Encode each symbol with 2 bits

alphabet symbol $x$	codeword $c(x)$
$A$	00
$C$	01
$G$	10
$T$	11

**Output of encoding:** the concatenation of the codewords for every symbol in the input sequence

### Example

- ▶ Input sequence:  $ACGTAA$
- ▶ Output:  $c(A)c(C)c(G)c(T)c(A)c(A) = 000110110000$
- ▶ Total length of encoding  $= 6 \cdot 2 = 12$  bits.

# Today

- 1 Recap
- 2 Data compression
- 3 Symbol codes and optimal lossless compression**
- 4 Prefix codes
- 5 Prefix codes and trees
- 6 The Huffman algorithm

# Symbol codes

**Symbol code:** a set of codewords where every input symbol is encoded **separately**.

**Symbol code:** a set of codewords where every input symbol is encoded **separately**.

Examples of symbol codes

- ▶  $C_0 = \{00, 01, 10, 11\}$  is a symbol code for  $\{A, C, G, T\}$ .
- ▶ ASCII encoding system: every character and special symbol on the computer keyboard is encoded by a different 7-bit binary string.

# Symbol codes

**Symbol code:** a set of codewords where every input symbol is encoded **separately**.

Examples of symbol codes

- ▶  $C_0 = \{00, 01, 10, 11\}$  is a symbol code for  $\{A, C, G, T\}$ .
- ▶ ASCII encoding system: every character and special symbol on the computer keyboard is encoded by a different 7-bit binary string.

## Remark 1.

$C_0$  and ASCII are *fixed-length* symbol codes: each codeword has the same length.



# Unique decodability

Decoding  $C_0$ ?

# Unique decodability

## Decoding $C_0$

- ▶ read 2 bits of the output;
- ▶ print the symbol corresponding to this codeword;
- ▶ continue with the next 2 bits.

## Decoding $C_0$

- ▶ read 2 bits of the output;
- ▶ print the symbol corresponding to this codeword;
- ▶ continue with the next 2 bits.

## Remark 2.

- ▶ *This decoding algorithm works for ASCII (replace 2 by 7)*
- ▶  $C_0$ , ASCII: *distinct input sequences have distinct encodings*

# Unique decodability

## Decoding $C_0$

- ▶ read 2 bits of the output;
- ▶ print the symbol corresponding to this codeword;
- ▶ continue with the next 2 bits.

## Remark 2.

- ▶ *This decoding algorithm works for ASCII (replace 2 by 7)*
- ▶  $C_0$ , ASCII: *distinct input sequences have distinct encodings*

## Definition 1.

A symbol code is **uniquely decodable** if, for any two **distinct** input sequences, their encodings are distinct.

# Lossless compression

- ▶ **Lossless compression:** compress and decompress without errors.
- ▶ Uniquely decodable codes allow for **lossless compression**.
- ▶ A symbol code achieves **optimal lossless compression** when it produces an encoding of **minimum length** for its input (among all uniquely decodable symbol codes).
- ▶ Huffman algorithm: provides a symbol code that achieves **optimal** lossless compression.

## Fixed-length vs variable-length codes

Chromosome map consists of 200 million bases as follows:

alphabet symbol $x$	frequency $freq(x)$
$A$	110 million
$C$	5 million
$G$	25 million
$T$	60 million

# Fixed-length vs variable-length codes

Chromosome map consists of 200 million bases as follows:

alphabet symbol $x$	frequency $freq(x)$
$A$	110 million
$C$	5 million
$G$	25 million
$T$	60 million

- ▶  $A$  appears much more often than the other symbols.
- ⇒ It might be best to encode  $A$  with fewer bits.
- ▶ Unlikely that the **fixed-length encoding**  $C_0$  is optimal

# Today

- 1 Recap
- 2 Data compression
- 3 Symbol codes and optimal lossless compression
- 4 Prefix codes**
- 5 Prefix codes and trees
- 6 The Huffman algorithm



# Prefix codes

Variable-length encodings

Code  $C_1$

alphabet symbol $x$	codeword $c(x)$
$A$	0
$C$	00
$G$	10
$T$	1

# Prefix codes

## Variable-length encodings

Code  $C_1$

alphabet symbol $x$	codeword $c(x)$
$A$	0
$C$	00
$G$	10
$T$	1

- $C_1$  is **not** unique decodable! E.g., 101110: *how to decode it?*

# Prefix codes

Variable-length encodings

Code  $C_2$

alphabet symbol $x$	codeword $c(x)$
$A$	0
$C$	110
$G$	111
$T$	10

# Prefix codes

## Variable-length encodings

Code  $C_2$

alphabet symbol $x$	codeword $c(x)$
$A$	0
$C$	110
$G$	111
$T$	10

- ▶  $C_2$  is uniquely decodable.
- ▶  $C_2$  is such that no codeword is a **prefix** of another.

# Prefix codes

Variable-length encodings

Code  $C_2$

alphabet symbol $x$	codeword $c(x)$
$A$	0
$C$	110
$G$	111
$T$	10

- ▶  $C_2$  is uniquely decodable.
- ▶  $C_2$  is such that no codeword is a **prefix** of another.

**Definition 2 (prefix codes).**

A symbol code is a **prefix code** if no codeword is a prefix of another.

# Decoding prefix codes

1. Scan the binary string from left to right until you've seen enough bits to match a codeword;
2. Output the symbol corresponding to this codeword.
  - ▶ Since no other codeword is a prefix of this codeword or contains it as a prefix, this sequence of bits cannot be used to encode any other symbol.
3. Continue starting from the next bit of the bit string.

# Decoding prefix codes

1. Scan the binary string from left to right until you've seen enough bits to match a codeword;
2. Output the symbol corresponding to this codeword.
  - ▶ Since no other codeword is a prefix of this codeword or contains it as a prefix, this sequence of bits cannot be used to encode any other symbol.
3. Continue starting from the next bit of the bit string.

Thus prefix codes allow for

- ▶ **unique decoding**;
- ▶ **fast decoding** (the end of a codeword is instantly recognizable).

# Decoding prefix codes

1. Scan the binary string from left to right until you've seen enough bits to match a codeword;
2. Output the symbol corresponding to this codeword.
  - ▶ Since no other codeword is a prefix of this codeword or contains it as a prefix, this sequence of bits cannot be used to encode any other symbol.
3. Continue starting from the next bit of the bit string.

Thus prefix codes allow for

- ▶ **unique decoding**;
- ▶ **fast decoding** (the end of a codeword is instantly recognizable).

Examples of prefix codes:  $C_0$ ,  $C_2$



# Prefix codes and optimal lossless compression

- ▶ Decoding a prefix code is very fast.
- ⇒ Would like to focus on prefix codes (rather than **all** uniquely decodable symbol codes) for achieving **optimal lossless compression**.
- ▶ Information theory guarantees that.
- ▶ So we can solely focus on prefix codes for optimal compression.

# Compression gains from variable-length prefix codes

Chromosome map: *do we gain anything by using  $C_2$  instead of  $C_0$  when compressing the map of 200 million bases?*

Input	
symbol $x$	$freq(x)$
$A$	110 million
$C$	5 million
$G$	25 million
$T$	60 million

Code $C_0$	
$x$	$c(x)$
$A$	00
$C$	01
$G$	10
$T$	11

Code $C_2$	
$x$	$c(x)$
$A$	0
$C$	110
$G$	111
$T$	10

# Compression gains from variable-length prefix codes

Chromosome map: *do we gain anything by using  $C_2$  instead of  $C_0$  when compressing the map of 200 million bases?*

Input		Code $C_0$		Code $C_2$	
symbol $x$	$freq(x)$	$x$	$c(x)$	$x$	$c(x)$
$A$	110 million	$A$	00	$A$	0
$C$	5 million	$C$	01	$C$	110
$G$	25 million	$G$	10	$G$	111
$T$	60 million	$T$	11	$T$	10

- ▶  $C_0$ : 2 bits  $\times$  200 million symbols = 400 million bits
- ▶  $C_2$ :  $1 \cdot 110 + 3 \cdot 5 + 3 \cdot 25 + 2 \cdot 60 = 320$  million bits
- ▶ Improvement of 20% in this example

# The optimal prefix code problem

## Input:

- ▶ Alphabet  $\mathcal{A} = \{a_1, \dots, a_n\}$
- ▶ Set  $P = \{p_1, \dots, p_n\}$  of probabilities over  $\mathcal{A}$  such that  $p_i = \Pr[a_i]$

**Output:** a binary prefix code  $C^* = \{c(a_1), c(a_2), \dots, c(a_n)\}$  for  $(\mathcal{A}, P)$ , where codeword  $c(a_i)$  has length  $\ell_i$  and the expected length of the code

$$L(C^*) = \sum_{a_i \in \mathcal{A}} p_i \cdot \ell_i$$

is **minimum** among all prefix codes.

# Example

## Chromosome example

Input	
symbol $x$	$\Pr(x)$
$A$	110/200
$C$	5/200
$G$	25/200
$T$	60/200

Code $C_0$	
$x$	$c(x)$
$A$	00
$C$	01
$G$	10
$T$	11

Code $C_2$	
$x$	$c(x)$
$A$	0
$C$	110
$G$	111
$T$	10

- ▶  $L(C_0) = 2$
- ▶  $L(C_2) = 1.6$
- ▶ *Coming up:*  $C_2$  is the output of the Huffman algorithm, hence an optimal encoding for  $(\mathcal{A}, P)$ .


# Today

- 1 Recap
- 2 Data compression
- 3 Symbol codes and optimal lossless compression
- 4 Prefix codes
- 5 Prefix codes and trees**
- 6 The Huffman algorithm

# Prefix codes and trees

- ▶ A **binary tree**  $T$  is a rooted tree such that each node that is not a leaf has at most two children.
- ▶ Binary tree for a prefix code: a branch to the left represents a 0 in the encoding and a branch to the right a 1.


Code $C_0$	
$x$	$c(x)$
$A$	00
$C$	01
$G$	10
$T$	11



# Prefix codes and trees

- ▶ A **binary tree**  $T$  is a rooted tree such that each node that is not a leaf has at most two children.
- ▶ Binary tree for a prefix code: a branch to the left represents a 0 in the encoding and a branch to the right a 1.

Code $C_2$	
$x$	$c(x)$
$A$	0
$C$	110
$G$	111
$T$	10





# Properties of binary trees representing prefix codes

1. *Where do alphabet symbols appear in the tree?*
2. *What do codewords correspond to in the tree?*
3. *Consider the tree corresponding to the optimal prefix code.  
Can it have internal nodes with one child?*

# Properties of binary trees representing prefix codes

1. Symbols must appear at the leaves of the tree  $T$  (*why?*)  
 $\Rightarrow T$  has  $n$  leaves.
2. Codewords  $c(a_i)$  are given by **root-to-leaf paths**.

Recall that  $\ell_i$  is the length of the codeword  $c(a_i)$  for input symbol  $a_i$ . Therefore, on the tree  $T$ ,  $\ell_i$  corresponds to the depth of  $a_i$  (we assume that the root is at depth 0).

$\Rightarrow$  Can rewrite the **expected length** of the prefix code as:

$$L(C) = \sum_{a_i \in \mathcal{A}} p_i \cdot \ell_i = \sum_{1 \leq i \leq n} p_i \cdot \text{depth}_T(a_i) = L(T).$$

3. **Optimal tree must be full**: all internal nodes must have exactly two children (*why?*).

## More on optimal tree

### Claim 1.

*There is an optimal prefix code, with corresponding tree  $T^*$ , in which the two lowest frequency characters are assigned to leaves that are siblings in  $T^*$  at maximum depth.*

# More on optimal tree

## Claim 1.


*There is an optimal prefix code, with corresponding tree  $T^*$ , in which the two lowest frequency characters are assigned to leaves that are siblings in  $T^*$  at maximum depth.*

## Proof.

By an **exchange argument**: start with a tree for an optimal prefix code and **transform** it into  $T^*$ . □

# Proof of Claim 1


- ▶ Let  $T$  be the tree for the optimal prefix code.
- ▶ Let  $\alpha, \beta$  be the two symbols with the smallest probabilities, that is,  $\Pr[\alpha] \leq \Pr[\beta] \leq \Pr[s]$  for all  $s \in \mathcal{A} - \{\alpha, \beta\}$ .
- ▶ Let  $x$  and  $y$  be the two siblings at maximum depth in  $T$ .



We want  $L(T) \geq L(T')$

# Proof of Claim 1

- ▶ Let  $T$  be the tree for the optimal prefix code.
- ▶ Let  $\alpha, \beta$  be the two symbols with the smallest probabilities, that is,  $\Pr[\alpha] \leq \Pr[\beta] \leq \Pr[s]$  for all  $s \in \mathcal{A} - \{\alpha, \beta\}$ .
- ▶ Let  $x$  and  $y$  be the two siblings at maximum depth in  $T$ .



We want  $L(T') \geq L(T'')$

## How do the expected lengths of the two trees compare?

$$\begin{aligned}L(T) - L(T') &= \sum_{a_i \in \mathcal{A}} \Pr[a_i] \cdot \text{depth}_T(a_i) - \sum_{a_i \in \mathcal{A}} \Pr[a_i] \cdot \text{depth}_{T'}(a_i) \\&= \Pr[\alpha] \cdot \text{depth}_T(\alpha) + \Pr[x] \cdot \text{depth}_T(x) \\&\quad - \Pr[\alpha] \cdot \text{depth}_{T'}(\alpha) - \Pr[x] \cdot \text{depth}_{T'}(x) \\&= \Pr[\alpha] \cdot \text{depth}_T(\alpha) + \Pr[x] \cdot \text{depth}_T(x) \\&\quad - \Pr[\alpha] \cdot \text{depth}_T(x) - \Pr[x] \cdot \text{depth}_T(\alpha) \\&= (\Pr[\alpha] - \Pr[x]) \cdot (\text{depth}_T(\alpha) - \text{depth}_T(x)) \geq 0\end{aligned}$$

- ▶ The third equality follows from the exchange.
- ▶ Similarly, exchanging  $\beta$  and  $y$  in  $T'$  yields  $L(T') - L(T'') \geq 0$ .
- ▶ Hence  $L(T) - L(T'') \geq 0$ .
- ▶ Since  $T$  is optimal, it must be  $L(T) = L(T'')$ .
- ▶ So  $T''$  is also optimal.

The claim follows by setting  $T^*$  to be  $T''$ .

# Building the optimal tree

Claim 1 tells us how to build the optimal tree **greedily!**

1. Find the two symbols with the lowest probabilities.
2. Remove them from the alphabet and replace them with a new **meta-character** with probability equal to the sum of their probabilities.
  - ▶ **Idea:** this meta-character will be the parent of the two deleted symbols in the tree.
3. Recursively construct the optimal tree using this process.

**Greedy algorithms:** *make a local (myopic) decision at every step that optimizes some criterion and eventually show that this is the optimal way for building the entire solution.*



# Today

- 1 Recap
- 2 Data compression
- 3 Symbol codes and optimal lossless compression
- 4 Prefix codes
- 5 Prefix codes and trees
- 6 The Huffman algorithm**

# Huffman algorithm

**Huffman**( $\mathcal{A}, P$ )

**if**  $|\mathcal{A}| = 2$  **then**

    Encode one symbol using 0 and the other using 1

**end if**

Let  $\alpha$  and  $\beta$  be the two symbols with the lowest probabilities

Let  $\nu$  be a new meta-character with probability  $\Pr[\alpha] + \Pr[\beta]$

Let  $\mathcal{A}_1 = \mathcal{A} - \{\alpha, \beta\} + \{\nu\}$

Let  $P_1$  be the new set of probabilities over  $\mathcal{A}_1$

$T_1 = \text{Huffman}(\mathcal{A}_1, P_1)$

**return**  $T$  as follows: replace leaf node  $\nu$  in  $T_1$  by an internal node, and add two children labelled  $\alpha$  and  $\beta$  below  $\nu$ .

## Remark 3.


*Output of **Huffman** procedure is a binary tree  $T$ ; the code for  $(\mathcal{A}, P)$  is its corresponding prefix code.*

# Example: recursive Huffman for chromosome map


Recursive call 1:  $\text{Huffman}(\{A, C, G, T\}, \{\frac{110}{200}, \frac{5}{200}, \frac{25}{200}, \frac{60}{200}\})$

Recursive call 2:  $\text{Huffman}(\{A, \nu_{CG}, T\}, \{\frac{110}{200}, \frac{30}{200}, \frac{60}{200}\})$


Recursive call 3:  $\text{Huffman}(\{A, \nu_{CGT}\}, \{\frac{110}{200}, \frac{90}{200}\})$



End of rec. call 3



End of rec. call 2



End of rec. call 1

**Proof:** by induction on the size of the alphabet  $n \geq 2$ .

- ▶ **Base case.** For  $n = 2$ , Huffman is optimal.
- ▶ **Hypothesis.** Assume that Huffman returns the optimal prefix code for alphabets of  $n$  symbols.
- ▶ **Induction Step.** Let  $\mathcal{A}$  be an alphabet of size  $n + 1$ ,  $P$  the corresponding set of probabilities.

Let  $T_1$  be the optimal (by the hypothesis) tree returned by our algorithm for  $(\mathcal{A}_1, P_1)$ , where  $\mathcal{A}_1, P_1, T_1$  as in the pseudocode. Let  $T$  be the final tree returned for  $(\mathcal{A}, P)$  by our algorithm. We claim that  $T$  is optimal.

We will prove the claim by contradiction; to this end, assume that  $T^*$  is the optimal tree for  $(\mathcal{A}, P)$  such that

$$L(T^*) < L(T). \tag{1}$$

# A useful fact

## Fact 3.

Let  $T$  be a binary tree representing a prefix code. If we replace siblings  $\alpha, \beta$  in  $T$  by a meta-character  $\nu$  where  $\Pr[\nu] = \Pr[\alpha] + \Pr[\beta]$ , we obtain a tree  $T_1$  such that

$$L(T) = L(T_1) + (\Pr[\alpha] + \Pr[\beta]).$$

## Proof.

**Notation:**  $d_T(a_i) = \text{depth}_T(a_i)$

- $\alpha, \beta$  are siblings in  $T$ , hence  $d_T(\alpha) = d_T(\beta)$ .
- $T$  differs from  $T_1$  only in that  $\alpha, \beta$  are replaced by  $\nu$ . Since  $d_{T_1}(\nu) = d_T(\alpha) - 1$ , we obtain

$$\begin{aligned} L(T) - L(T_1) &= \Pr[\alpha]d_T(\alpha) + \Pr[\beta]d_T(\beta) - (\Pr[\alpha] + \Pr[\beta])d_{T_1}(\nu) \\ &= \Pr[\alpha] + \Pr[\beta]. \end{aligned} \tag{2}$$



## Correctness (cont'd)

- ▶ Claim 1 guarantees there is such an optimal tree for  $(\mathcal{A}, P)$  where  $\alpha, \beta$  appear as siblings at maximum depth.
- ▶ W.l.o.g. assume that  $T^*$  is such an optimal tree. By Fact 3, if we replace siblings  $\alpha, \beta$  in  $T^*$  by  $\nu'$  where  $\Pr[\nu'] = \Pr[\alpha] + \Pr[\beta]$ , the resulting tree  $T_1^*$  satisfies  $L(T^*) = L(T_1^*) + (\Pr[\alpha] + \Pr[\beta])$ .
- ▶ Similarly, the tree  $T$  returned by the Huffman algorithm satisfies  $L(T) = L(T_1) + (\Pr[\alpha] + \Pr[\beta])$ .
- ▶ By the induction hypothesis, we have  $L(T_1^*) \geq L(T_1)$  since  $T_1$  is optimal for alphabets of size  $n$ . Hence

$$L(T^*) = L(T_1^*) + \Pr[\alpha] + \Pr[\beta] \geq L(T_1) + \Pr[\alpha] + \Pr[\beta] = L(T), \quad (3)$$

where the inequality follows from the induction hypothesis.

- ▶ Equation (3) contradicts assumption 1. Thus  $T$  must be optimal.

# Implementation and running time

1. Straightforward implementation:  $O(n^2)$  time
2. Store the alphabet symbols in a priority queue implemented as a **binary min-heap** with **keys** their probabilities
  - ▶ **Operations:** Initialize ( $O(n)$ ), Extract-min ( $O(\log n)$ ), Insert ( $O(\log n)$ )

Total time:  $O(n \log n)$  time

For an iterative implementation of Huffman, see your textbook.

# Example: iterative Huffman for chromosome map




Input ( $\mathcal{A}, P$ )

symbol $x$	$\text{Pr}(x)$
$A$	110/200
$C$	5/200
$G$	25/200
$T$	60/200

$\rightarrow$

Output code

symbol $x$	$c(x)$
$A$	0
$C$	110
$G$	111
$T$	10





# Beyond Huffman coding

- ▶ Huffman algorithm provides an optimal **symbol** code.
- ▶ Codes that encode larger blocks of input symbols might achieve better compression.
- ▶ Storage on noisy media: *what if a bit of the output of the compressor is flipped?*
  - ▶ Decompression cannot carry through.
  - ▶ Need **error correction** on top of compression.